Introduction to RAG

- **Retrieval-Augmented Generation (RAG)** is an AI framework that enhances LLMs with external knowledge retrieval
- Combines the strengths of **retrieval-based** and **generation-based** AI approaches
- Introduced by Meta AI Research in 2020 as a method to improve factual accuracy
- Enables LLMs to access and leverage information beyond their training data
- Produces more accurate, up-to-date, and contextually relevant responses

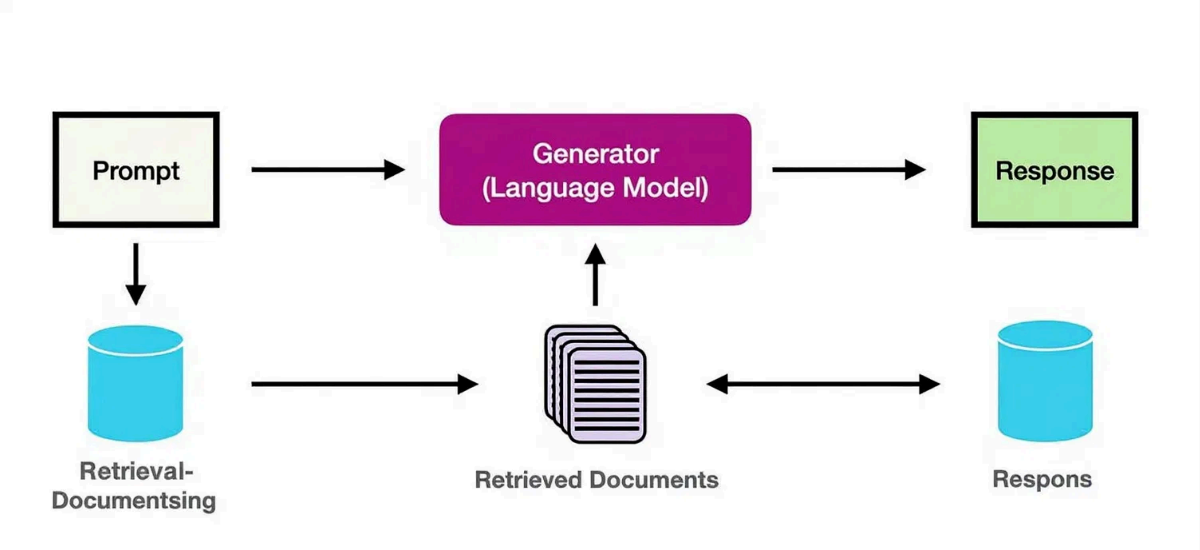

Limitations of Standard LLMs

- **Hallucinations:** LLMs can generate plausible but factually incorrect information
- **Knowledge Cutoff:** Limited to information available up to their training cutoff date
- **Domain Limitations:** Lack specialized knowledge in specific fields or proprietary information
- **Source Attribution:** Difficulty in citing sources or providing evidence for generated content
- **Context Window:** Limited ability to process and reference large amounts of information

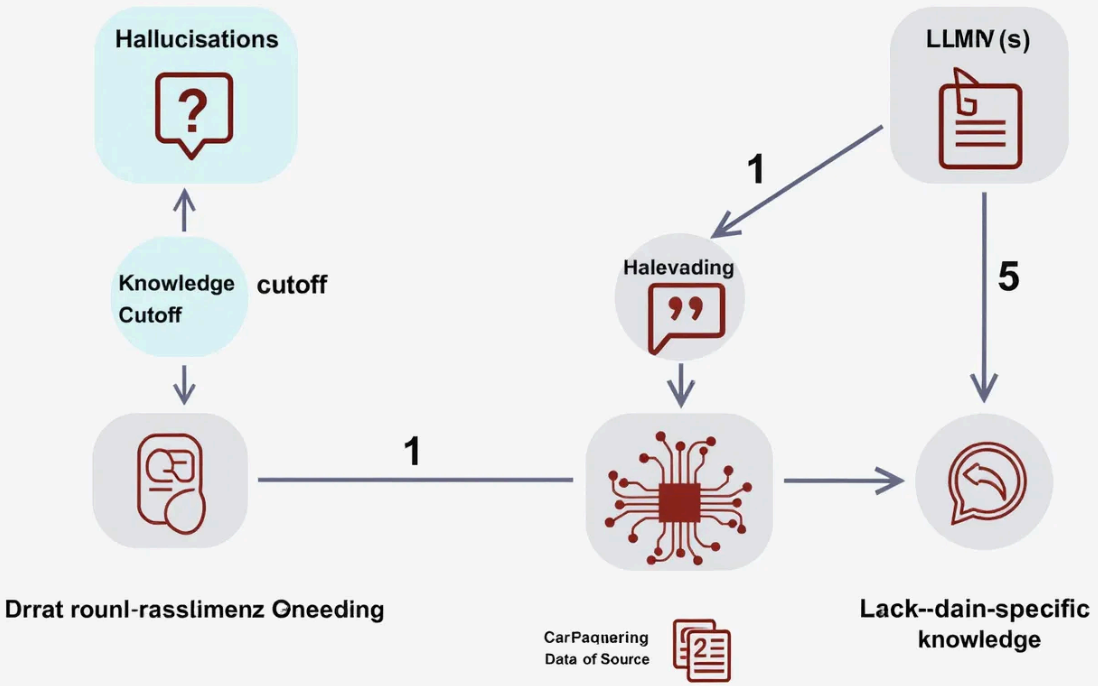

How RAG Works

- **1. Retriever:** Searches for relevant information from external knowledge sources
- **2. Knowledge Base:** Contains indexed documents, often using vector embeddings for semantic search
- **3. Context Augmentation:** Retrieved information is added to the prompt as context
- **4. Generator:** LLM uses the augmented context to produce accurate, grounded responses
- **5. End-to-End Flow:** Query → Retrieval → Augmentation → Generation → Response

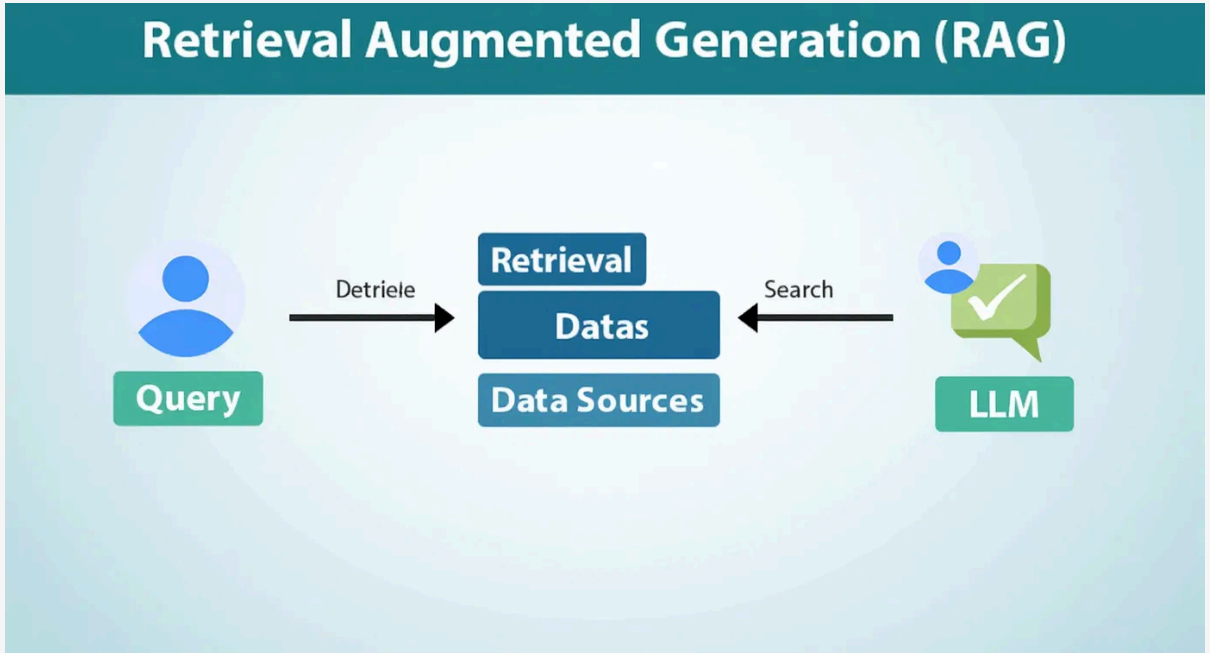

Key Components and Architecture

- **Knowledge Base:** Document store containing indexed information (PDFs, web pages, databases)
- **Embedding Model:** Converts text into vector representations for semantic search
- **Retrieval Pipeline:** Query processing, vector search, and relevance ranking
- **Integration Layer:** Combines retrieved context with user query for the LLM
- **Generator:** LLM that produces the final response using augmented context

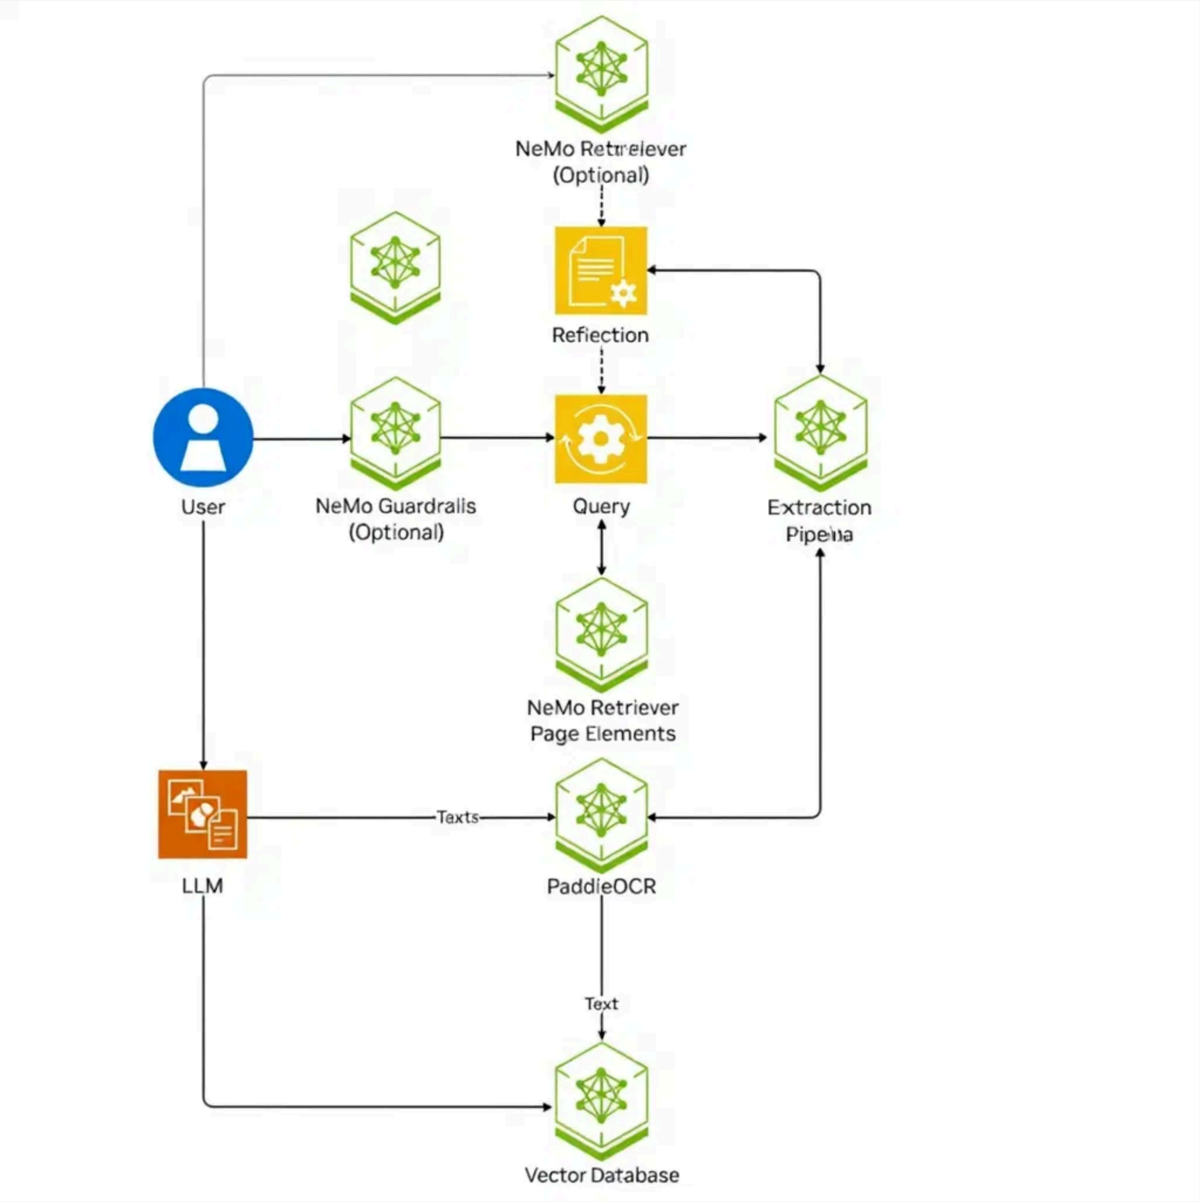

Real-World Applications

- **Customer Support:** Chatbots that access company knowledge bases to provide accurate, contextual responses
- **Healthcare:** Clinical decision support systems that retrieve medical literature and patient records
- **Legal Research:** Document analysis tools that search case law and legal precedents
- **Content Creation:** Writing assistants that incorporate factual information and citations
- **Education:** Personalized tutoring systems that access textbooks and learning materials

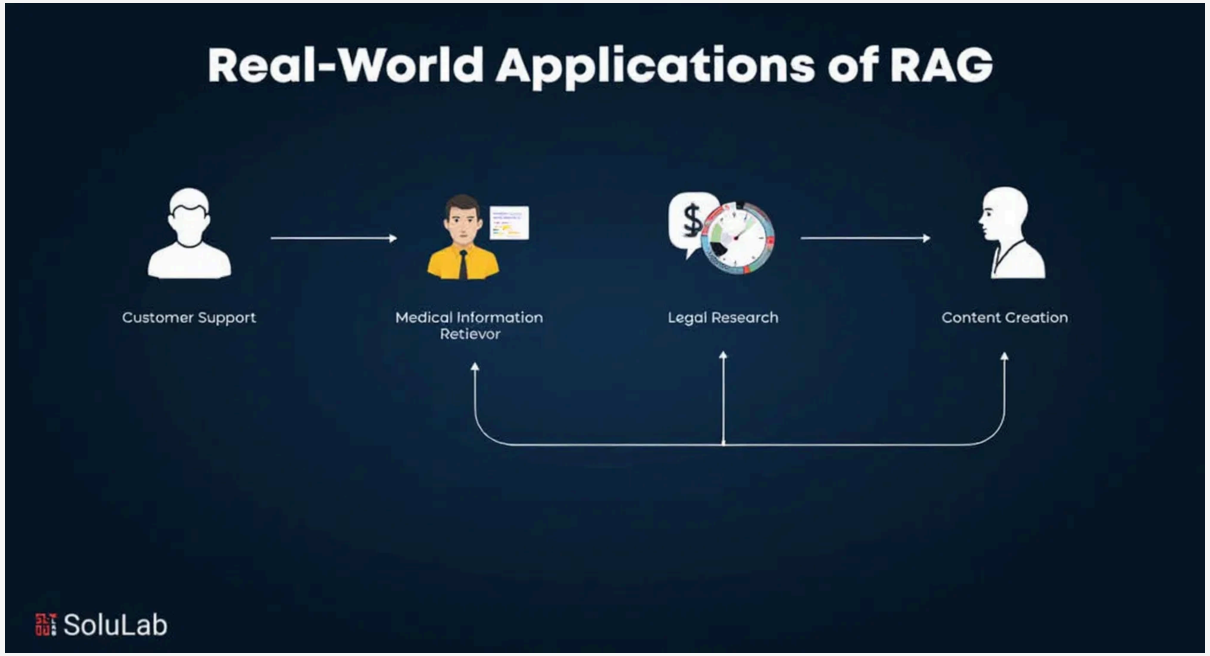

Benefits and Challenges

Benefits

- **Improved Accuracy:** Reduces hallucinations by grounding responses in factual data
- **Current Information:** Accesses up-to-date knowledge beyond training cutoff
- **Domain Expertise:** Incorporates specialized knowledge from proprietary sources

Challenges

- **System Complexity:** More components to maintain and optimize
- **Retrieval Quality:** Results depend on the quality of search and indexing
- **Computational Cost:** Additional processing overhead compared to standard LLMs